PRUTHVI KADAM

(862) 230-9988 • pruthvivkadam@gmail.com • linkedin.com/in/pruthvi-kadam • Newark, NJ

PROFESSIONAL SUMMARY

Results-driven Data Scientist and Machine Learning Engineer with 2+ years of professional experience building scalable data pipelines, deploying ML models, and delivering actionable insights through advanced analytics. Proficient in Python, SQL, Scikit-learn, and NLP, with a strong academic foundation in Data Science (GPA 3.75, NJIT). Adept at translating complex data into business value through visualization tools such as Power BI and Tableau. Seeking entry-to-mid-level Data Scientist, Data Analyst, or ML Engineer roles where I can drive data-driven decision-making.

EDUCATION

M.S. in Data Science *Sep 2024 – May 2026*
New Jersey Institute of Technology, Newark, NJ | *GPA: 3.75 / 4.0*

B.S. in Data Science *Aug 2019 – May 2022*
Thakur College of Science and Commerce, Mumbai, India | *GPA: 9.72 / 10.0*

TECHNICAL SKILLS

Languages: Python, R, SQL, C#, Java
ML & Data Science: Machine Learning, Deep Learning, NLP, Feature Engineering, A/B Testing, Statistical Analysis, Model Deployment, Big Data Analytics, LLMs
Libraries & Frameworks: Pandas, NumPy, Scikit-learn, Matplotlib, Seaborn, Hugging Face Transformers, vLLM, Flask, Tree-sitter
Visualization & BI: Tableau, Power BI, Matplotlib
Tools & Platforms: AWS EMR, Snowflake, Jupyter, Git, MySQL, Postman, Microsoft Excel, Visual Studio
Certifications: Bloomberg Market Concepts | Microsoft Power Automate

PROFESSIONAL EXPERIENCE

Data Engineer & Software Developer | **ACTY Systems India Pvt. Ltd.** *Mumbai, India* *Jun 2022 – Aug 2024*
- Engineered and maintained ETL pipelines using Python and PySpark to ingest, transform, and load 500K+ employee HR records, reducing data processing latency by 35% through query optimization on Snowflake.
- Automated 10+ repetitive data workflows using Python scripting and .NET-based automation tools, eliminating ~15 hours of manual effort per week and improving team throughput by 40%.
- Applied data augmentation and statistical validation techniques across multiple software release cycles, reducing system downtime by 20% and improving overall data integrity.
- Designed and deployed a plagiarism-detection system for coding submissions using ML similarity algorithms (cosine similarity, TF-IDF), achieving 92%+ detection accuracy and ensuring fair evaluation for 500+ participants.
- Collaborated cross-functionally with product, QA, and engineering teams to integrate analytics-driven features into the Employee Management System, improving HR reporting efficiency and leadership decision support.
- Mentored 5+ junior developers on data engineering best practices and testing workflows, accelerating team onboarding time and improving code review efficiency by 150%.

PROJECTS

C# Instruction-Tuning Pipeline for LLM Fine-Tuning *Jan 2025*

Designed an end-to-end data pipeline to fine-tune StarCoder2-15B on C# tasks by adapting the Self-OSS-Instruct framework; curated 1,942 high-quality code samples from The Stack v2 corpus using SHA-1 hashing for deduplication.
- Applied Tree-sitter AST parsing to structurally filter syntactically valid C# functions, ensuring dataset quality for downstream LLM training and instruction-response alignment.
- Orchestrated vLLM-powered inference to generate natural-language instruction-response pairs at scale, integrating automated self-validation heuristics that filtered ~30% of low-quality outputs.
- **Tech stack:** Python, vLLM, Hugging Face Transformers, Tree-sitter, StarCoder2-15B, hashlib

Student Performance Analyzer — *New Jersey Institute of Technology* ***Sep 2024***

- Built an end-to-end academic analytics tool using Python and Pandas to identify at-risk students by analyzing GPA trends, attendance, and assignment completion rates across a dataset of 1,000+ student records.
- Developed interactive visualizations using Matplotlib and Seaborn to surface performance bottlenecks, enabling faculty to take targeted interventions; deployed via Flask as a real-time web reporting dashboard.
- **Tech stack:** Python, Pandas, Matplotlib, Seaborn, Flask

Mumbai Housing Price Prediction — *Thakur College of Science and Commerce* ***Aug 2019***

- Developed a regression-based ML model using Scikit-learn (Random Forest, XGBoost) to predict Mumbai real estate prices with 85% accuracy, leveraging features such as location, square footage, amenities, and proximity to transit.
- Performed end-to-end data preprocessing with Pandas including outlier removal, missing value imputation, and feature engineering on a 20,000+ row housing dataset, improving model performance by 12% over baseline.
- Deployed the predictive model as a Flask web application enabling real-time property price estimation with an interactive user interface.
- **Tech stack:** Python, Scikit-learn, XGBoost, Pandas, Flask